Replace with title page

CITS4009 Computational Data Analysis
Semester Two 2021

This Paper Contains: **5** pages **(including title page)**

Time allowed: 2 hours

INSTRUCTIONS:

This paper contains 6 questions.

TOTAL: 60 MARKS

Students should attempt ALL questions.

Answers are to be written in the answer booklet provided.
Question paper is to be collected with the answer booklet.

- Answers should be concise rather than lengthy.
- If you think that a question is ambiguous, state clearly any assumptions that you make in constructing your answer.
- For questions that require you to write R code, minor syntactic errors will not be penalised; however, syntactic errors that obscure the meaning of your answer might cost you marks. Pseudo code may be given partial marks.

Students can bring in one sheet of A4-size paper with hand-written or typed notes on both sides.

This page has been left intentionally blank

1. a) (2 marks) Briefly state the differences between *hypothesis generation* and *hypothesis confirmation*.

   b) (2 marks) Use an example to illustrate that *Exploratory Data Analysis* is an iterative process.

   c) (2 marks) What are the three essential components of the layered grammar of graphics that *gglot* implements? Give an example for each component.

   d) (2 marks) What visualisation (or plot) is most suitable to illustrate the covariation between two continuous variables? Give an example and write R code using the *ggplot2* library to illustrate your answer.

   e) (2 marks) Briefly explain the differences between a *histogram* plot and a *density* plot. When would it be more suitable to use a density plot than a histogram plot?

2. a) (4 marks) Given below is a data frame `df` showing the *for sale* prices of some properties in a good suburb in Western Australia. The property type and the number of bedrooms are in the first two columns. The last column contains the prices in 10^3 dollars.

Type	Num.bedrooms	Price
Villa	2	525
House	3	1200
Apartment	1	460
Apartment	2	950
House	3	1000
Villa	1	395
House	4	1300
Villa	2	600

   i. (2 marks) Describe a plot that is suitable for visualising variable `Type` versus variable `Num.bedrooms`. Write R code using the *ggplot2* library to illustrate your answer.

   ii. (2 marks) Describe a plot that is suitable for visualising variable `Type` versus variable `Price`. Write R code using the *ggplot2* library to illustrate your answer.

   b) (1 mark) Describe when it would be suitable to convert a continuous variable into a categorical one.

   c) (2 marks) Referring to the data frame `df` in part a) above, suppose that we want to convert the `Price` column to the following levels to form a new categorical column called `Price.Range`:

   - `Low` if Price $\leq$ 500;
   - `Medium` if 500 $<$ Price $\leq$ 1,000;
   - `High` if Price $>$ 1,000.

   Write R code to add `Price.Range` to the data frame.

   d) (1 mark) Explain what is meant by *listwise deletion* in data cleaning.

   e) (2 marks) Describe two different ways for imputing missing values in a numerical column. mean/ ratio, machine learning predict result

3. a) (2 marks) Explain what *z-normalisation* is. Is it suitable for detecting outliers? Explain your answer.

   b) (4 marks) For a given vector `v`, five numbers are output by `boxplot.stats(v)$stats`. Explain what each of these numbers represents. By inspecting these numbers alone, can we determine whether `v` is free of outliers? Explain your answer.

   c) (4 marks) Given two data frames `authors` and `books`, which have a common `surname` column, as shown below:

surname	nationality	deceased
Tukey	US	yes
Venables	Australia	yes
Ripley	NZ	no
Tierney	US	no
Winton	UK	no

surname	title	other.author
Tukey	Exploratory Data Analysis	NA
Venables	Modern Applied Statistics	Ripley
Tierney	LISP-STAT	NA
Ripley	Spatial Statistics	NA
Ripley	Stochastic Simulation	NA
McNeil	Interactive Data Analysis	NA
R Core	An Introduction to R	Venables & Smith

   i. (3 marks) Explain the difference between the *inner join* and *left outer join* operations on these data frames.

   ii. (1 marks) Write R code to show how the output tables can be produced from the *inner join* and *left outer join* operations on `authors` and `books`.

4. Each row of the data frame `df` below shows the measurement of a type of blood test and whether the patient currently smokes (`smoke`), has never smoked (`never`), or has smoked before but has now quit (`quit`). The last column of the data frame is a binary variable indicating whether the patients have been diagnosed with a type of cancer.

Patient	Smoke	Test	Cancer
1	never	0.56	negative
2	quit	1.10	positive
3	smoke	1.50	positive
4	never	1.20	negative
5	smoke	1.60	positive
6	quit	0.98	negative

   a) (4 marks) Explain how a *decision tree* classifier partitions the data frame and assigns a piece-wise constant to each partition. You can use any input feature as the first variable. The output variable that the classifier should predict is the `Cancer` column.

   b) (3 marks) Explain how the *k-nearest neighbours* classifier works for this data frame for predicting the `Cancer` variable. List two aspects that need to be considered during data preparation for this classifier.

   c) (3 marks) Explain how the *receiver operating characteristic curve* and the double density plots can be used to compare the performance of the two binary classifiers above.

5. a) (3 marks) Define what a typical Null model would be like for the data frame in Question 4 above, where the response variable that we want to predict is the `Cancer` column. Write R code to show the predicted probability produced by your Null model.

   b) (2 marks) Explain how the `dist()` function in R can be used to find the distances between data points.

   c) (3 marks) What is the *k-means* algorithm designed for? Outline the steps involved in this algorithm.

   d) (2 marks) Explain what the *Calinski-Harabasz index* measures.

6. Given below are the first 8 observations of data frame `df` for a simple *dry bean* dataset. It has 3 classes in the last column and 4 features (or variables): `Perimeter`, `roundedness`, `ShapeFactor1`, and `ShapeFactor2`.

Perimeter	roundness	ShapeFactor1	ShapeFactor2	Class
954.496	0.864	0.00598	0.00119	Cali
716.507	0.954	0.00641	0.00250	Seker
1040.323	0.853	0.00561	0.00105	Cali
776.180	0.877	0.00632	0.00224	Seker
898.660	0.698	0.00605	0.00224	Seker
941.694	0.855	0.00593	0.00132	Barbunya
1105.912	0.851	0.00514	0.00108	Cali
750.314	0.945	0.00609	0.00247	Seker

   a) (1 mark) Write R code to relabel the `Class` column to the following values:

   - 1, to replace `Seker`, and
   - 0, to replace `Cali` and `Barbunya`

   for binary classification.

   b) (2 marks) Write an R function called `calDeviance`, which should take in two arguments, `ytrue` (for the ground truth vector) and `ypred` (for the predicted vector). The function should compute the *deviance* and return it as the output value. You may assume that the saturated model has zero deviance.

   c) (2 marks) Write R code to split the dataset into a training set and a calibration set. Use an 80/20 ratio for the splitting.

   d) (5 marks) For each of the 4 features, write R code to

      i. (2 marks) train a *logistic regression* classifier model using the training set (Hint: you can use the `glm` function),

      ii. (2 marks) apply the trained model on the calibration set, and

      iii. (1 mark) call the `calDeviance` function above and print the output value.