



# Lab 03 - EDA for Single Variable

## Learning outcomes

In this lab you will learn and practise

- Customise your UI in Shiny
- Exploratory data analysis for a single variable
- Understanding the stats in histograms, `boxplot`, and `barplot`.
- Data frame subsetting, functions and loops

## Part A: UI design

In this part, we will explore some UI design in Shiny.

**Step 1:** Begin by creating a Shiny app file. Go to File → New File → Shiny Web App. You'll encounter two application types: single file or multiple files. While both are functional, for this example, we'll use the single file option. Enter a name for your application, such as "week4App", and then click on the "create" button. This action will generate an "app.R" file.

**Step 2:** Inside the "app.R", you'll find a sample code. To visualise the application, you can simply click on "Run App" and interact with the provided sample.

**Step 3:** Next, let's customise the "app.R" file according to your preferences. You can begin by clearing the existing sample code and load dataset by using:

```
df <- read.csv("https://archive.ics.uci.edu/ml/machine-learning-databases/wine")
```

Note that, `df` is just a variable name.

```

library(ggplot2)
library(shiny)

# load dataset
df <- read.csv("https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.csv")
# Define UI for application e.g. draws a histogram
ui <- fluidPage(
  titlePanel("JUST A NAME"),
  sidebarLayout(
    sidebarPanel(
      # Add sidebar content here
    ),
    mainPanel(
      plotOutput(outputId = "show")
    )
  )
)
# Define server logic required draw a histogram
server <- function(input, output) {
  output$show <- renderPlot({
    # Add server logic here
  })
}
# Run the application
shinyApp(ui = ui, server = server)

```

**Step 4:** Let's explore different input options within the `sidebarPanel`.

1. **Slider Input:** Begin by incorporating the provided code snippet into the `sidebarPanel`. When you run or reload the app, you will observe a slider element displayed on the web interface.

```

sliderInput(inputId = "bins", label = "Number of bins:", min=1, max=50, value=10)

```

2. **Select Input:** Modify your `sidebarPanel` according to the given instructions below. Then, reload the app to witness the changes.

```

sidebarPanel(
  sliderInput(inputId = "bins", label = "Number of bins:", min=1, max=50, value=10),
  selectInput(inputId = "x1", label = "Choose X:", choices = names(df)),
  selectInput(inputId = "x2", label = "Choose a quality:", choices = c(df$quality, df$pH)),
  selectInput(inputId = "x3", label = "Choose from quality and pH:", choices = c(df$quality, df$pH))
),

```

3. **Radio Button:** Integrate the following code snippet into your `sidebarPanel` and reload the app.

```
radioButtons(inputId = "animal", label = "What's your favourite animal?", choi
```

You can explore more input options in <https://mastering-shiny.org/basic-ui.html#inputs>.

**Step 5 (optional):** Step 4 generally covers the essentials of basic design fulfillment. However, for a more personalised touch and highly customised Shiny apps, you can incorporate HTML elements into your UI. For further information on various HTML elements, you can refer to [https://www.w3schools.com/html/html\\_elements.asp](https://www.w3schools.com/html/html_elements.asp).

```


library(shiny)
df <- read.csv("https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.csv")
# Define UI for application that draws a histogram
ui <- fluidPage(
  h1("This is a heading"),
  p("This is some text", class = "my-class1"),
  tags$ul(
    tags$li("First bullet"),
    tags$li("Second bullet")
  ),
  # Application title
  titlePanel("JUST A NAME"),
  h2("heading 2"),
  sidebarLayout(
    sidebarPanel(
      sliderInput(inputId = "bins", label = "Number of bins:", min=1, max=10),
      selectInput(inputId = "x1", label = "Choose X:", choices = names(df)),
      selectInput(inputId = "x2", label = "Choose a quality:", choices = names(df[, 3:13])),
      selectInput(inputId = "x3", label = "Choose from quality and pH:", choices = names(df[, 3:13])),
      radioButtons(inputId = "animal", label = "What's your favourite animal?", choices = c("cat", "dog", "bird", "fish")),
      # adding the new div tag to the sidebar
      tags$div(class="header", checked=NA,
        tags$p("Ready to take the Shiny tutorial? If so"),
        tags$a(href="https://shiny.posit.co/r/articles/build/html/creating-a-shiny-app.html")
      ),
      h3("This is also a heading"),
      p("This is also some text", class = "my-class2"),
      tags$ul(
        tags$li("First bullet 2"),
        tags$li("Second bullet 2")
      ),
    ),
    mainPanel(
      plotOutput("distPlot")
    )
  )
)
# Define server logic required to draw a histogram
server <- function(input, output) {
  output$distPlot <- renderPlot({
    # Draw a histogram
  })
}
# Run the application
shinyApp(ui = ui, server = server)

```

**Step 6:** Having acquired the skills to craft a customised UI, let's shift our focus to the `mainPanel`. The function `plotOutput` plays a pivotal role in Shiny. When employing

`plotOutput`, it's important to pair it with `renderPlot` in your server section. More functions are provided below.

## Outputs `render*()` and `*Output()` functions work together to add R output to the UI.




**DT::renderDataTable**(expr, options, searchDelay, callback, escape, env, quoted, outputArgs)

**dataTableOutput**(outputId)



**renderImage**(expr, env, quoted, deleteFile, outputArgs)

**imageOutput**(outputId, width, height, click, dblclick, hover, brush, inline)



**renderPlot**(expr, width, height, res, ..., alt, env, quoted, execOnResize, outputArgs)

**plotOutput**(outputId, width, height, click, dblclick, hover, brush, inline)

`"data.frame" 3 obs. of 2 variables:  
$ Sepal.Length: num 5.1 5.8 5.7  
$ Petal.Length : num 5.1 5.1 5.2`

**renderPrint**(expr, env, quoted, width, outputArgs)

**verbatimTextOutput**(outputId, placeholder)

	MeanSepal.Length	MeanPetal.Length	MeanSpecies
1	5.10	5.10	Setosa
2	5.80	5.80	Versicolour
3	5.70	5.70	Versicolour

**renderTable**(expr, striped, hover, bordered, spacing, width, align, rownames, colnames, digits, na, ..., env, quoted, outputArgs)

**tableOutput**(outputId)

foo

**renderText**(expr, env, quoted, outputArgs, sep)

**textOutput**(outputId, container, inline)



**renderUI**(expr, env, quoted, outputArgs)

**uiOutput**(outputId, inline, container, ...)  
**htmlOutput**(outputId, inline, container, ...)

These are the core output types. See [htmlwidgets.org](https://htmlwidgets.org) for many more options.

`"distPlot"` is the outputId. Feel free to replace it with any suitable name. Remember, if you modify the outputId, the corresponding adjustment should also be made in the `renderPlot` function within the server section. For instance, if I've defined the id as `plotOutput("hist")`, then the server's `renderPlot` should align accordingly:

```
output$hist <- renderPlot({  
  })
```

## Part B: Server design

Let's design the server logic. To illustrate, I'll use the `geom_histogram()` function. Assuming we're working with an R dataframe named `df` and focusing on the column `fixed.acidity`, here's the code for plotting a histogram:

```
ggplot(df, aes(x = fixed.acidity)) +geom_histogram(bins = 20, color = "white")
```

To adapt this function for use within the Shiny server, follow these steps:

```
output$hist <- renderPlot({  
  ggplot(df, aes_string(x = input$x1)) +geom_histogram(bins = 20, color = "white")  
})
```


In this code:

- `input$x1` refers to the selected choice from the first `selectInput` where you've defined `inputId = "x1"`.
- We've replaced `aes` with `aes_string` to accommodate dynamic input variable names.

## Part C: EDA for Single Variable

Note that you **do not need** to perform this part in Shiny. Let's continue analysing the Wine data from last week, we read the data into a data frame and did a simple histogram plot for the `fixed.acidity`. For example, using the following code:

```
wineData <- read.csv("https://archive.ics.uci.edu/ml/machine-learning-database  
fixedAcidity <- hist(wineData$fixed.acidity)
```



If it doesn't work, then download the `winequality-red.csv` file and read it from your local disk.

Inspect the variable `fixedAcidity` that stores the output from the `hist()` function above. In particular, look at `fixedAcidity$breaks`, `fixedAcidity$mids` and `fixedAcidity$counts`. What are the lengths of these three vectors?

**Write some R code to understand how the stats are calculated in the R built-in Histogram**

**Exercise 1:** Write an R statement that involves using `fixedAcidity$counts` and `fixedAcidity$mids` to find out how many observations falling into bin `(7,8]` in the above histogram. Here, `(` indicates *non-inclusion*, and `]` indicates *inclusion* (i.e., the number of observations whose *fixed acidity* values are greater than 7 but less than or equal to 8). Your R code should produce the following:

```
## [1] 496
```

**Exercise 2:** Extend Exercise 1 above as follows: Write a `for` loop which would count the numbers of observations whose fixed acidity values fall into the 12 bins `(4,5]`, `(5,6]`, ..., `(15,16]`. Save these numbers of observations to a vector named `counts`. Your `counts` vector should have the following values:

```
## [1] 9 62 291 496 300 188 118 76 39 12 3 5
```

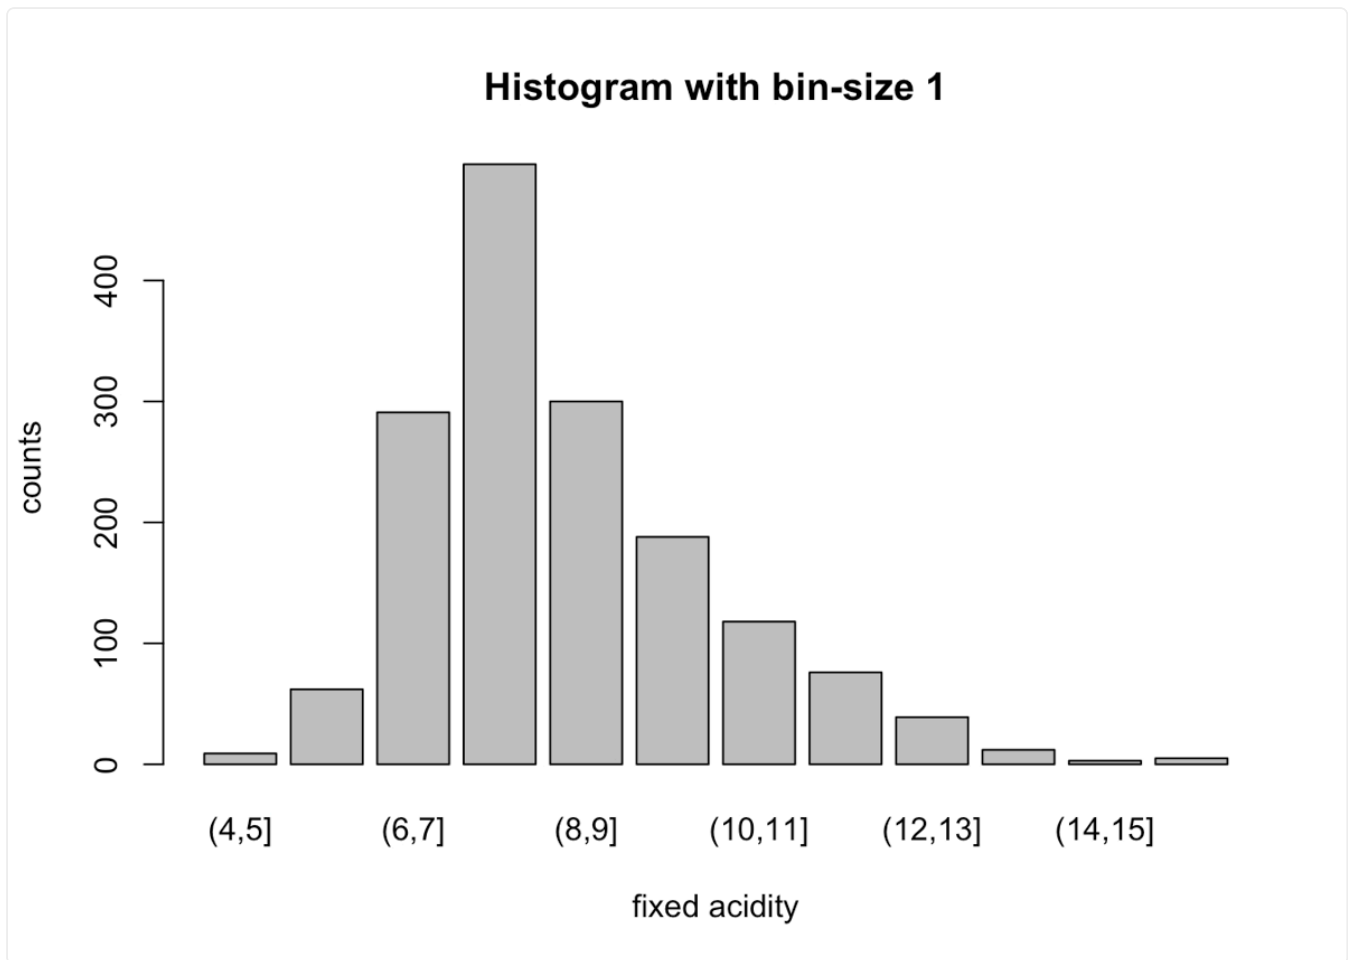
**Hint:** You can use the following block of code to create a zero vector of length  $n$  (in our case, we want  $n=12$ ):

```
count <- c()
n <- 12
for (i in 1:n) count <- c(count, 0)
```

Alternatively and more efficiently, since the `counts` vector is going to store integer values only, we can simply do:

```
counts <- integer(12)
```

Plot your `counts` vector using the R built-in `barplot()` function. Set some appropriate parameters so that your plot has a meaningful title, axis labels and tick marks, like the diagram shown below:



Note that R automatically omits some tick marks to avoid over-crowding the axis in the plot. To get the label below each histogram bin as shown in the diagram above, try

```
bins <- seq(4,15)
name_str <- paste("(", bins, ",", bins+1, "]", sep="")
```


and use the variable `name_str` appropriately when constructing your barplot.

Compare your plot with the histogram from `hist()`.

**Excercise 3:** Take a look at values on the vertical axis in the following R built-in Density Histogram, when `freq` is set to `FALSE`.

```
hist(wineData$fixed.acidity, freq=FALSE)
```

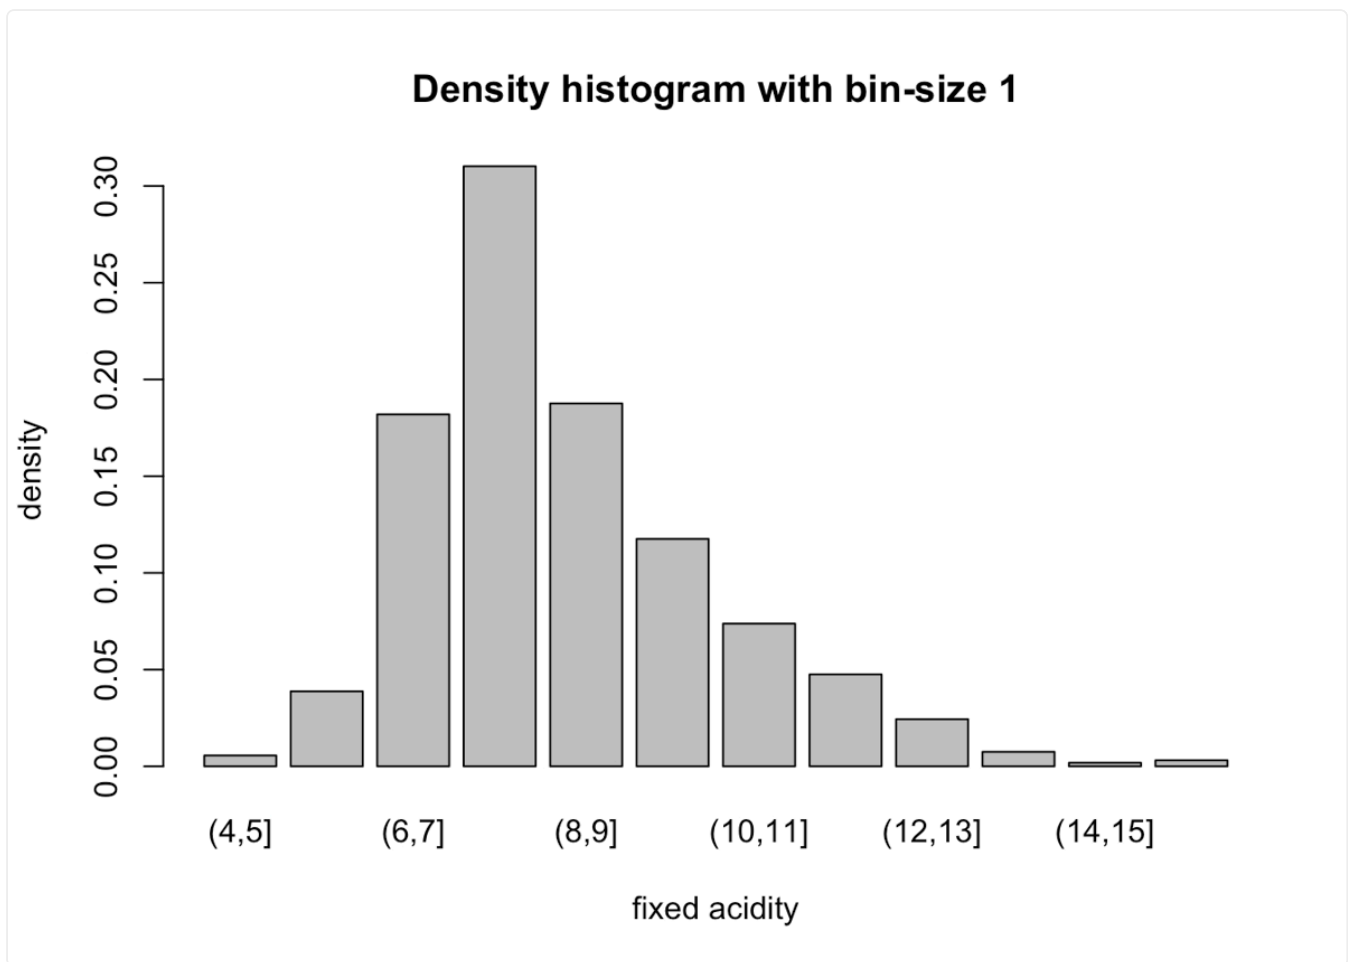





Reproduce the values along the *density* axis for each bin based on your `counts` vector that you worked out earlier:

```
freqs <- counts / sum(counts)
```

and plot these values using the built-in `barplot()` function. Your plot should look like this:




**Exercise 4:** Plot the histogram and density diagrams using the `geom_histogram` and `geom_density` functions from the *ggplot2* library into one figure. Use `binwidth = 1` for the histogram. You can make `geom_histogram` to plot the density instead of raw counts by using `mapping = aes(x = fixed.acidity, y = ..density..)`, where `..density..` is `ggplot` calculated stat.



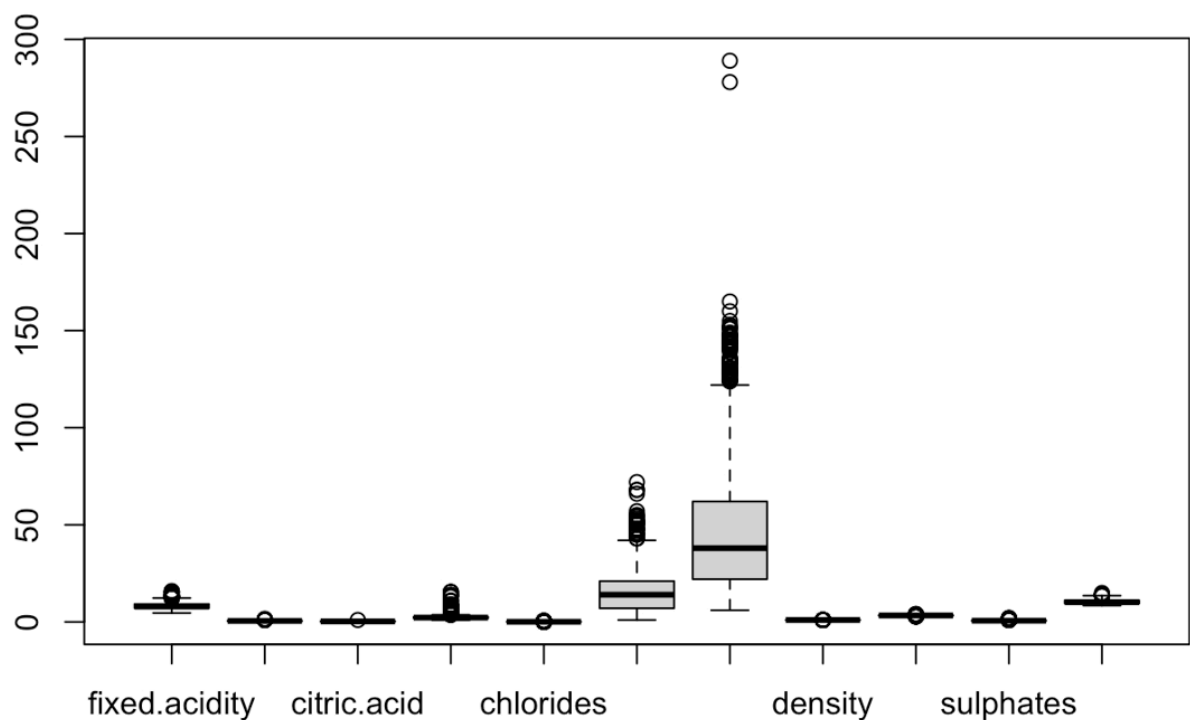
## Box plots

**Exercise 1:** Use box plot (either the basic `boxplot()` function or the `geom_boxplot()` from the *ggplot2* library to plot the `wineData$fixed.acidity` variable. You should get a diagram that looks like this (if you use `geom_boxplot()` from *ggplot2*):




What does this plot tell you?

**Exercise 2:** Write R code to produce the box plots of all the variables (except for the `quality` column,) side-by-side in the same figure, like the diagram show below (if you use the basic `boxplot()` ):



To produce the above side-by-side box plot using `ggplot` and flip the coordinate axes using `coord_flip()`, you can consider the code below, where `stack(wineData)` creates a new two-column data frame from `wineData` by stacking the 12 columns of `wineData` row-wise. The column names of this new data frame are `ind` and `values`. In the code below, we rename the axes to *variable* and *value*.

```
ggplot(stack(wineData)) +
  geom_boxplot(mapping = aes(x = ind, y = values)) +
  labs(x='variable', y='value') +
  coord_flip()
```



Try saving `stack(wineData)` to a variable `df` and inspect `df`. How many rows are there in the data frame `df`? Verify that `nrow(df)` is equal to `nrow(wineData) * ncol(wineData)`.

**Exercise 3:** Read the help page for `boxplot.stats` and write a line of R code to obtain the five value stats for `fixed.acidity`. You should get a vector that looks like this:

```
## [1] 4.6 7.1 7.9 9.2 12.3
```

**Exercise 4:** Try to find the median using your own understanding rather than using the build-in functions. **Hint:** sort the vector of values in `fixed.acidity` and find the index that splits the sorted values in half. Try to work out the code yourself before looking at the solution below.

```
median_index <- floor(length(wineData$fixed.acidity)/2) + 1
sorted <- sort(wineData$fixed.acidity)
cat("The median value of fixed.acidity is", sorted[median_index])
```

```
## The median value of fixed.acidity is 7.9
```

**Optional exercise 1:** Put your R code into a function called `myMedian` that would find and return the median of any vector of numerical values. So

`myMedian(wineData$fixed.acidity)` should give the following (same as `median(wineData$fixed.acidity)`):

```
print(myMedian(wineData$fixed.acidity))
## [1] 7.9
```

**Optional exercise 2:** Modify your `myMedian` function to have an optional argument `q`. Let's call this new function `myQuantile`. It should return the appropriate quantile value of the input vector. For example, if the input vector is `wineData$fixed.acidity`, then the lower quartile (when `q = 0.25`) or upper quartile (when `q = 0.75`) should output the following:

```
cat("The lower quartile is", myQuantile(wineData$fixed.acidity, q=0.25))
## The lower quartile is 7.1
cat("The upper quartile is", myQuantile(wineData$fixed.acidity, q=0.75))
## The upper quartile is 9.2
```

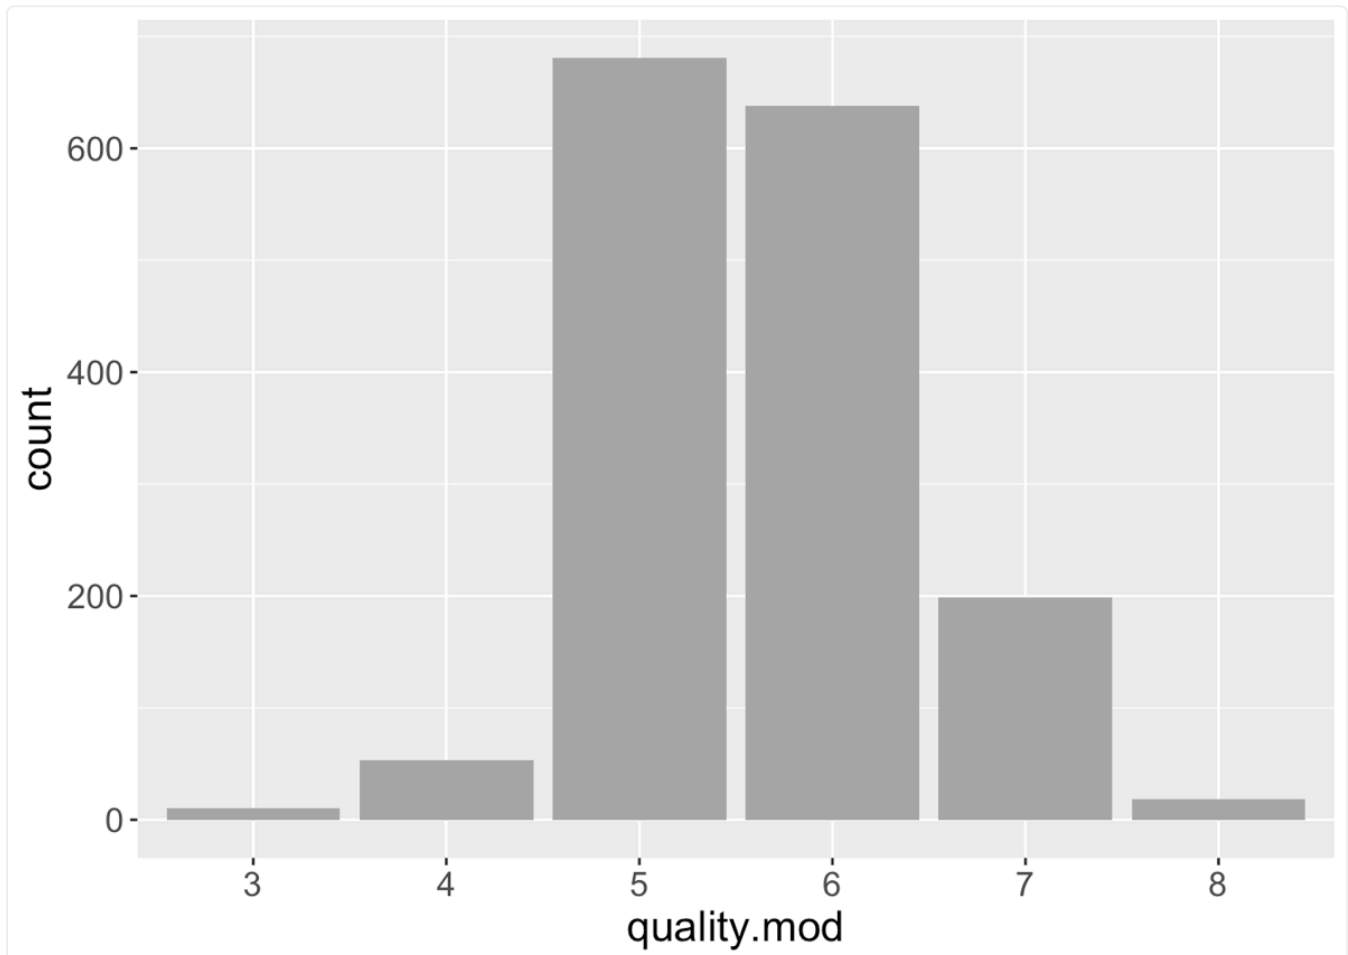
**Optional exercise 3:** Write a function called `computeQuantiles` which should take in a data frame (e.g., variable `wineData` above) and, depending on the value of the optional argument `q`, output a vector of numbers containing the lower quartile, median, or upper quartile values. The length of the vector should be equal to the number of numeric columns in the data frame. **Hint:** Use `ncol()` to find out the number of variables (columns) in the data frame, `is.numeric()` to find out if a given column contains numerical values. For instance, the lower quartiles and upper quartiles of all the numerical columns of the *red wine* data (variable `wineData`) returned by the function should be

```
print(computeQuantiles(wineData, q=0.25))
## [1] 7.1000 0.3900 0.0900 1.9000 0.0700 7.0000 22.0000 0.9956 3.2100
## [10] 0.5500 9.5000 5.0000
print(computeQuantiles(wineData, q=0.75))
## [1] 9.20000 0.64000 0.42000 2.60000 0.09000 21.00000 62.00000 0.9978
## [9] 3.40000 0.73000 11.10000 6.00000
```

## Bar chart

Use `str()` to look at the type of each column of the imported *red wine* data. The last column `quality` is of type integer, which should really be factors.

- Extract the `quality` column, turn it into factors, and store the results as a new column called `quality.mod` in the data frame.
- Use `ggplot`'s `geom_bar()` to plot the distribution of `quality.mod`. Repeat this operation for the distribution of `quality`. How different are the two plots? Repeat this operation again for another numerical variable, such as `fixed.acidity`, and see how the plot looks like.



## Continue the learning journey on Swirl

Type

```
library(swirl)
swirl()
```

and complete lessons 1-6 and the **Exploratory Data Analysis** course. You may need to uninstall the R Programming course. Inside `swirl`, type

```
uninstall_course("R Programming")
```

`swirl()` can get tricky at times, if you run into problem, please contact the Lab Facilitator.



# Try out R Markdown

Try out the R Markdown (.Rmd) if you haven't done so. You need to compose a R Markdown file for your EDA for project 1.

- **Cheat Sheet:** <https://www.rstudio.com/wp-content/uploads/2015/02/rmarkdown-cheatsheet.pdf>

Previous  
Markdown Basics

Next  
Solution to Lab 03 Part C

Last updated 1 month ago